

Breaking the Feedback Trap: Understanding and Stabilizing Recurrent Feedback Learning in Medical Image Segmentation

Anonymous Authors

Abstract—Recurrent feedback architectures in medical vision promise iterative boundary refinement by re-injecting past prediction masks into intermediate feature representations across training epochs. However, these networks frequently suffer from chronic, runaway false-positive over-segmentation. Crucially, while modern asymmetric boundary loss functions (e.g., Tversky loss) effectively suppress over-segmentation in standard feedforward backbones, their optimization benefit is completely neutralized once recurrent feedback is engaged.

In this work, we present a foundational mechanistic study of recurrent gradient dynamics, diagnosing the structural root cause of this failure: the *Feedback Trap*. We discover that the core vulnerability lies in the dual role forced upon recurrent predictions: using prediction history simultaneously as forward spatial guidance and as an unconstrained backward optimization path causes catastrophic error amplification. In conventional implementations, hard binary gating yields zero gradients almost everywhere for internal attention layers ($\frac{\partial \text{keep}}{\partial \text{mask}} \equiv 0$), while unchecked recurrent gradients induce severe Batch Normalization drift ($D_{\text{KL}} = 33.63$) that permanently locks the encoder into hallucinated background noise.

To resolve this pathology, we formulate the *Feedback Firewall* principle: forward spatial guidance must be strictly decoupled from backward optimization. We realize this via *Detached Soft-OR Gating*—a zero-parameter formulation combining continuous Boolean relaxation with the stop-gradient operator ($\text{detach}(m_{\text{fg}})$). This guarantees continuous gradient flow to internal attention proportional to background uncertainty ($\frac{\partial \text{keep}}{\partial \text{mask}} = 1 - m_{\text{fg}}$) while completely severing recursive error propagation ($\frac{\partial \text{keep}}{\partial m_{\text{fg}}} \equiv 0$).

Through extensive multi-seed benchmarks on Kvasir-SEG (Sessile), our approach slashes False Positive Rate by 42.8% (from 4.30% down to 2.46%), elevates mean Dice by +4.77 pp (to 29.95%), and reduces inter-seed variance by 66.5% (paired Wilcoxon $p < 0.001$). A 3×2 factorial analysis empirically proves that *Detached Soft-OR* successfully breaks the trap, reactivating asymmetric loss optimization. Furthermore, a 4-way ablation strictly disentangles the benefits of continuous relaxation from gradient detachment. Finally, external zero-shot cross-center evaluation on CVC-ClinicDB ($N = 612$, $p = 1.61 \times 10^{-15}$) and architectural universality verification on Recurrent Residual U-Net (R2U-Net) confirm that our Gradient Decoupling principle provides a universal foundation for robust recurrent learning in biomedical computer vision.

Index Terms—Recurrent Feedback, Gradient Decoupling, Feedback Firewall, Medical Image Segmentation,

Feedback Trap, Detached Soft-OR, Mechanistic Deep Learning

I. INTRODUCTION

Colorectal cancer (CRC) represents one of the leading causes of cancer-related mortality worldwide, with early endoscopic detection and resection of precancerous polyps serving as the clinical gold standard for prevention [1], [2]. Among varied morphology types, **sessile and flat polyps** (Paris Classification Types IIa and IIb) present the highest diagnostic hazard during colonoscopy [3]. Due to their low height profile, irregular borders, and textural indistinguishability from surrounding healthy mucosa, these subtle lesions are frequently missed or inaccurately delineated by automated segmentation algorithms [4], [5].

To address these subtle boundary ambiguities, deep learning architectures have increasingly adopted **recurrent feedback mechanisms** [6]–[8]. A prominent exemplar is the Feature Attention Network (FANet) [6], which introduces cross-epoch recurrent feedback. In recurrent vision networks, intermediate or past prediction masks (m) generated in preceding passes are re-injected as auxiliary attention priors to early encoder feature representations. The theoretical premise is compelling: by iteratively processing intermediate hypotheses through spatial gating modules (MixPool), the network should progressively sharpen ambiguous contours and prune false positives through multi-pass self-guidance [6], [9].

A. The Empirical Paradox: Over-Segmentation and Loss Neutralization

Despite theoretical promise, deploying recurrent feedback architectures in biomedical segmentation reveals a severe, paradoxical failure mode: **chronic false-positive over-segmentation**. In clinical colonoscopy cohorts, error audits reveal that over 84% of recurrent network errors are concentrated in false positives—bleeding predictions far into non-lesion mucosa (extending up to 40.7 pixels beyond the true histological boundary).

Standard machine learning intuition suggests rectifying false-positive bias via **asymmetric loss functions** (such as the Tversky loss [10] or asymmetric Focal loss [11]), which assign disproportionately higher penalties to false positives ($\alpha > \beta$). Indeed, when evaluated in a purely feedforward setting (with recurrent feedback disconnected), asymmetric

Tversky loss delivers an immediate, statistically significant reduction in false positives (-3.25 pp, $p = 0.005$).

However, our extensive factorial experiments uncover a baffling empirical paradox: **the moment recurrent feedback is engaged, the beneficial effect of asymmetric loss is completely abolished** (exhibiting a detrimental interaction effect of $+4.17$ pp). Rather than eliminating spurious boundaries, the recurrent feedback network traps the model in an echo chamber of its own past mistakes, rendering modern loss engineering entirely ineffective.

B. The Discovery: The Dual-Role Pathology (The Feedback Trap)

Through systematic representational and gradient diagnostics, we identify the exact mechanical pathology causing this phenomenon, which we term the **Feedback Trap**. The core failure stems from a pervasive design flaw in recurrent feedback learning: **forcing prediction history to act simultaneously as forward spatial guidance and as an unconstrained backward optimization path**.

In conventional implementations, the spatial feature gating operator combines the learnable internal attention map fmask with the downsampled recurrent mask m_{fg} via a hard thresholding operator:

$$\text{keep}_{\text{hard}} = \max(\mathbb{I}(\text{fmask} > 0.5), m_{\text{fg}}). \quad (1)$$

This formulation triggers two catastrophic mathematical failures:

- 1) **Gradient Paralysis of Internal Attention:** Because the indicator function $\mathbb{I}(\cdot > 0.5)$ has a derivative of zero almost everywhere, $\frac{\partial \text{keep}}{\partial \text{fmask}} \equiv 0$. The learnable convolutional layers dedicated to extracting internal spatial attention receive zero direct supervision from the segmentation loss objective.
- 2) **Toxic Recurrent Backpropagation and Error Amplification:** Concurrently, non-zero gradients backpropagate freely through the recurrent mask branch (m_{fg}). When the network produces an early false-positive hallucination, that erroneous mask is reinjected in the subsequent iteration. Because the network backpropagates through this unvalidated recursive loop, early encoder Batch Normalization distributions drift catastrophically ($D_{\text{KL}} = 33.63$ at `e1.r1.bn3`), and mask gradient norms collapse by 79.6%. The network optimizes its representations to validate its own past predictions rather than ground-truth anatomy, permanently overpowering any loss-level penalty.

C. Framing: A Mechanistic Study of Recurrent Gradient Decoupling

Disentangling Mechanistic Principles from SOTA Chasing: We explicitly frame this paper as a **mechanistic deep learning study** rather than an empirical effort to chase leaderboards. Our explicit scientific goal is to isolate, diagnose, and resolve the dual-role structural flaw

of recurrent feedback in biomedical vision. We formulate the foundational **Feedback Firewall** principle:

Forward spatial guidance must be strictly decoupled from backward gradient optimization.

By enforcing gradient detachment on recurrent feedback tensors ($\text{detach}(m_{\text{fg}})$), the recurrent prediction functions purely as an uncorrupted forward spatial prior, completely preventing recursive error amplification.

Justifying Adversarial Stress-Testing: To stress-test this pathology, we benchmark our method on the Kvasir-SEG Sessile subset. Flat, poorly demarcated sessile lesions represent the ultimate adversarial environment where the Feedback Trap—characterized by runaway false-positive over-segmentation—is most destructive. By curing the trap in this ambiguous regime, we demonstrate the fundamental robustness of the architectural principle.

D. Contributions of This Work

To break the Feedback Trap without discarding the intrinsic benefits of recurrent self-guidance, we propose **Detached Soft-OR Gating**—a principled formulation combining continuous Boolean relaxation with the Feedback Firewall. Our contributions are threefold:

- 1) **Discovery and Diagnosis of the Feedback Trap:** We provide the first systematic diagnosis of gradient paralysis and loss neutralization in recurrent medical segmentation networks, proving that unconstrained recurrent backpropagation creates an echo chamber of amplified errors documented via layer-wise Batch-Norm drift, feature cosine similarity, and gradient tracking.
- 2) **The Feedback Firewall and Detached Soft-OR Formulation:** We formulate the Gradient Decoupling principle and realize it via Detached Soft-OR gating with zero learnable parameters ($\Delta\Theta = 0$). We prove analytically that this guarantees continuous gradient flow strictly proportional to background uncertainty ($\frac{\partial \text{keep}}{\partial \text{fmask}} = 1 - m_{\text{fg}}$) while strictly severing degenerate recurrent error loops ($\frac{\partial \text{keep}}{\partial m_{\text{fg}}} \equiv 0$).
- 3) **Empirical Disentanglement, Generalization, and Universality:** Through a comprehensive 3×2 factorial analysis, we empirically prove that our method reactivates asymmetric loss optimization. A 4-way ablation study strictly disentangles the benefits of continuous relaxation from gradient detachment. On Kvasir-SEG (Sessile) across 5 seeds, our approach cuts False Positive Rate by 42.8% ($4.30\% \rightarrow 2.46\%$), elevates mean Dice by $+4.77$ pp ($p < 0.001$), and reduces variance by 66.5%. Crucially, external zero-shot cross-center evaluation on CVC-ClinicDB ($N = 612$, $p = 1.61 \times 10^{-15}$) proves robust out-of-distribution transfer, while verification on Recurrent Residual U-Net (R2U-Net) confirms that our Feedback Firewall is a universal stabilization principle for recurrent vision architectures.

II. RELATED WORK

A. Medical Image and Polyp Segmentation

Automated lesion segmentation from colonoscopy video frames plays a vital role in computer-aided detection (CADE) and diagnosis (CADx) systems [1], [2]. Beginning with the foundational U-Net architecture [12] and its multi-scale variants such as UNet++ [13] and ResUNet [14], deep convolutional networks have established standard benchmarks in biomedical segmentation. To better capture subtle boundary transitions, attention-guided models have been proposed, including Attention U-Net [15], PraNet [16] (which utilizes parallel reverse attention to model polyp contours), and HarDNet-MSEG [17]. More recently, vision transformers such as TransUNet [18], Swin-Unet [19], and Polyp-PVT [20] have leveraged self-attention to model long-range spatial dependencies.

However, standard feedforward architectures remain inherently single-pass: intermediate representations cannot be iteratively re-examined or refined in light of downstream contextual decisions. In sessile polyp segmentation, where lesions are flat and visual margins are ambiguous, single-pass feedforward models often struggle to balance sensitivity against excessive background false alarms.

B. Feedback and Recurrent Refinement Networks

In biological vision, anatomical feedback projections from higher cortical areas to early visual processing stages outnumber feedforward projections by a significant margin [21], [22]. Inspired by this principle, recurrent feedback networks have been developed to simulate iterative cognitive refinement in artificial neural networks [7], [23]. In general computer vision, Feedback Networks [23] and recurrent convolutional networks (RCNN) [24] demonstrate that re-routing higher-level predictions to lower-level feature extractors improves object localization and contextual disambiguation.

In medical imaging, Recurrent U-Net [8] and particularly the Feature Attention Network (FANet) [6] introduced cross-epoch feedback loops. FANet introduces the **MixPool** module, reinjecting the previous epoch’s binary prediction mask directly into the encoder and decoder stages to guide feature extraction.

The Overlooked Research Gap: Existing literature in recurrent segmentation operates on the heuristic assumption that reinjecting prediction masks into intermediate layers automatically encourages iterative refinement. Crucially, these studies treat the recurrent coupling as a black box, completely overlooking the gradient dynamics of the feedback interface. In this work, we demonstrate that hard-thresholded mask injection creates a catastrophic *Feedback Trap*: it attenuates gradient backpropagation by over 79%, paralyzes internal attention learning, and causes early encoder representations to drift uncontrollably.

C. Loss Formulations for Imbalanced Medical Segmentation

Segmentation of small, irregular lesions is heavily impacted by class imbalance, where background pixels vastly

outnumber foreground lesion pixels [10]. Standard Cross-Entropy often biases predictions toward the background, prompting the widespread adoption of the Dice loss [25] and soft Jaccard loss [26].

To explicitly penalize false-positive over-segmentation or false-negative misses, asymmetric and boundary-aware loss formulations have emerged. The Tversky loss [10] generalizes the Dice index by introducing weighting parameters α and β to independently penalize false positives and false negatives. The Focal Tversky loss [27] further modulates hard examples, while Asymmetric Unified Focal loss [11] adaptively balances precision and recall.

Our Architectural Perspective: Contemporary research overwhelmingly attempts to suppress false positives by designing increasingly intricate loss functions. However, our factorial experiments reveal that such loss engineering is entirely neutralized when paired with conventional recurrent feedback loops. Rather than proposing another loss formulation, our work addresses the structural root of the problem: by redesigning the recurrent gating mechanism to eliminate zero gradients and detach correlative feedback paths, we **unleash the latent corrective power of existing asymmetric loss functions**, enabling them to suppress over-segmentation effectively.

D. Gradient Routing and Stop-Gradient Dynamics in Deep Learning

In deep representation learning, the stop-gradient operator ($\text{detach}(\cdot)$) has played a pivotal role in preventing degenerate representational collapse, most notably in Siamese self-supervised learning (e.g., SimSiam [28], BYOL [29]). In these frameworks, severing the backward computation graph along one pathway prevents trivial constant solutions and asymmetric drift without requiring negative sampling or complex architectures. Similarly, in reinforcement learning and actor-critic methods, target networks and detached value estimators are essential to stabilize recursive temporal-difference updates [30]. Despite its transformative impact in these domains, the principle of **gradient decoupling** has been largely overlooked in recurrent biomedical vision networks, where recurrent predictions have historically been permitted to flow unconstrained into the backward computation graph. Our work establishes the first rigorous theoretical and empirical application of the Feedback Firewall in recurrent medical vision.

III. METHODOLOGY: GRADIENT DECOUPLING AND THE FEEDBACK FIREWALL

In this section, we dissect the internal gradient mechanics of recurrent feature aggregation, prove the mathematical paralysis inherent in conventional coupled gating, and formulate our proposed **Feedback Firewall** and **Detached Soft-OR** architecture.

A. The Classical MixPool: Architecture and The Hard-Gating Flaw

The Feature Attention Network incorporates recurrent refinement across four encoder stages (e_1, e_2, e_3, e_4) and four decoder stages (d_1, d_2, d_3, d_4). At each level, intermediate feature tensors $x \in \mathbb{R}^{B \times C \times H \times W}$ and the recurrent feedback mask $m \in \mathbb{R}^{B \times 1 \times H_{in} \times W_{in}}$ are processed through a MixPool block.

Within MixPool, input feature x is first routed to a lightweight convolutional attention sub-network $\mathcal{F}_{att}$, generating a continuous single-channel attention prior:

$$\text{fmask} = \sigma\left(\text{Conv}_{1 \times 1}\left(\text{ReLU}\left(\text{BN}\left(\text{Conv}_{3 \times 3}(x)\right)\right)\right)\right) \in [0, 1]^{B \times 1 \times H \times W}. \quad (2)$$

Simultaneously, the feedback mask m is spatially downsampled via max pooling to match the spatial resolution (H, W) of the current layer, yielding $m_{fg} \in [0, 1]^{B \times 1 \times H \times W}$.

In the original formulation [6], these two spatial gates are fused into a single binary selection mask $\text{keep}_{\text{hard}}$ via a hard thresholding operator:

$$\text{keep}_{\text{hard}} = \max(\mathbb{I}(\text{fmask} > 0.5), m_{fg}). \quad (3)$$

The gated features are then split across dual convolutional pathways and recombined:

$$x_1 = \text{Conv}_1(x \odot \text{keep}_{\text{hard}}), \quad x_2 = \text{Conv}_2(x), \quad (4)$$

$$\text{MixPool}(x, m) = [x_1 \parallel x_2] \in \mathbb{R}^{B \times C \times H \times W}. \quad (5)$$

1) *The Calculus of Failure: Zero-Gradient and Toxic Coupling:* Let $\mathcal{L}$ denote the scalar segmentation objective. Applying the chain rule to backpropagate gradients from x_1 back to the parameters θ_{att} of the attention generator $\mathcal{F}_{att}$:

$$\frac{\partial \mathcal{L}}{\partial \theta_{att}} = \frac{\partial \mathcal{L}}{\partial x_1} \cdot \frac{\partial x_1}{\partial \text{keep}_{\text{hard}}} \cdot \frac{\partial \text{keep}_{\text{hard}}}{\partial \text{fmask}} \cdot \frac{\partial \text{fmask}}{\partial \theta_{att}}. \quad (6)$$

Because $\text{keep}_{\text{hard}}$ utilizes the indicator function $\mathbb{I}(\text{fmask} > 0.5)$, the distributional derivative of the Heaviside step function is the Dirac delta:

$$\frac{\partial}{\partial u} \mathbb{I}(u > 0.5) = \delta(u - 0.5) = 0 \quad \forall u \neq 0.5. \quad (7)$$

Consequently, the third factor in Eq. (6) evaluates to:

$$\frac{\partial \text{keep}_{\text{hard}}}{\partial \text{fmask}} \equiv 0 \quad \text{almost everywhere.} \quad (8)$$

As a direct consequence, **zero gradient reaches the parameters θ_{att} of the attention generator**. The network cannot adapt its internal spatial feature selector in response to task-specific segmentation errors.

Simultaneously, the recurrent branch receives backward gradients through automatic differentiation:

$$\frac{\partial \mathcal{L}}{\partial m_{fg}} = \frac{\partial \mathcal{L}}{\partial x_1} \cdot \frac{\partial x_1}{\partial \text{keep}_{\text{hard}}} \cdot \frac{\partial \text{keep}_{\text{hard}}}{\partial m_{fg}} \neq 0. \quad (9)$$

When the network produces an early false-positive prediction, this erroneous mask is reinjected into the encoder. Because gradients flow through m_{fg} back into the recursive

history without internal attention guidance, the encoder parameters are updated to maximize consistency with past errors rather than anatomical boundaries. This causes the internal feature representations of healthy background mucosa to collapse into lesion-like embeddings.

B. Proposed Solution: Detached Soft-OR Gating

To eradicate the Feedback Trap while preserving the inductive bias of iterative multi-pass refinement, we propose **Detached Soft-OR Gating**. Our method redesigns the gating operator along two dimensions:

1) *Continuous Boolean Relaxation via Soft-OR:* Instead of discontinuous hard thresholding, we relax the logical disjunction ($A \vee B$) into continuous probabilistic space. In probability theory, the union of two independent events A and B is given by:

$$\begin{aligned} P(A \cup B) &= P(A) + P(B) - P(A \cap B) \\ &= 1 - (1 - P(A))(1 - P(B)). \end{aligned} \quad (10)$$

Applying this formulation directly to the continuous attention activations $a \in [0, 1]$ and downsampled feedback mask $m_{fg} \in [0, 1]$:

$$\begin{aligned} \text{keep}_{\text{soft}} &= a + m_{fg} - (a \odot m_{fg}) \\ &= 1 - (1 - a) \odot (1 - m_{fg}), \end{aligned} \quad (11)$$

where $\odot$ denotes element-wise Hadamard multiplication.

2) *Severing the Degenerate Loop via Gradient Detachment:* To prevent erroneous recurrent priors from corrupting early encoder layers via recursive backpropagation, we enforce the stop-gradient operator on the recurrent mask tensor inside all MixPool modules:

$$\text{keep}_{\text{detached}} = a + \text{detach}(m_{fg}) - a \odot \text{detach}(m_{fg}). \quad (12)$$

C. Theoretical Properties and Analytical Guarantees

Our proposed Detached Soft-OR formulation provides three analytical guarantees:

Theorem 1 (Continuous Gradient Flow with Dynamic Uncertainty Scaling). *Under the Detached Soft-OR formulation (Eq. 12), the gradient of the feature gate with respect to the internal attention map fmask is strictly non-zero, continuous, and dynamically proportional to background uncertainty:*

$$\frac{\partial \text{keep}_{\text{detached}}}{\partial \text{fmask}} = 1 - m_{fg} \geq 0. \quad (13)$$

Proof. Taking the partial derivative of Eq. (12) with respect to $a \equiv \text{fmask}$:

$$\begin{aligned} \frac{\partial \text{keep}_{\text{detached}}}{\partial a} &= \frac{\partial}{\partial a} [a + \text{detach}(m_{fg}) - a \cdot \text{detach}(m_{fg})] \\ &= 1 - \text{detach}(m_{fg}). \end{aligned} \quad (14)$$

Because $\text{detach}(\cdot)$ behaves as an identity in forward evaluation and zero in backward differentiation, the effective backward derivative is $1 - m_{fg}$. Since $m_{fg} \in [0, 1]$, this quantity is strictly non-negative and bounded in $[0, 1]$. $\square$

Clinical Intuition of the Derivative:

- When $m_{fg} \rightarrow 1$ (high-confidence past lesion interior), $\frac{\partial keep}{\partial fmask} \rightarrow 0$. The internal attention branch does not waste gradient capacity refining already-confident lesion cores.
- When $m_{fg} \rightarrow 0$ (ambiguous boundary or healthy background mucosa), $\frac{\partial keep}{\partial fmask} \rightarrow 1$. The internal attention branch receives maximum gradient magnitude, empowering the network to learn precise boundary discrimination and prune false positives.

Theorem 2 (Complete Severance of Degenerate Error Propagation). *Under Detached Soft-OR, backward automatic differentiation through the recurrent mask branch is identically zero:*

$$\frac{\partial keep_{detached}}{\partial m_{fg}} \equiv 0. \quad (15)$$

Proof. Direct consequence of applying the stop-gradient operator $\text{detach}(\cdot)$ on m_{fg} . $\square$

Zero Parameter Overhead: The entire transformation in Eq. (12) requires exactly **zero learnable parameters** ($\Delta\Theta = 0$) and introduces zero additional FLOPs beyond elementary tensor arithmetic.

IV. EXPERIMENTS AND EMPIRICAL RESULTS

To evaluate the empirical validity of our theoretical analysis, we conduct comprehensive experiments across two distinct clinical colonoscopy benchmarks and external medical vision domains:

- 1) **Kvasir-SEG (Sessile):** A specialized cohort of 200 sessile and flat colorectal lesions (Paris IIa/IIb) from Oslo University Hospital [4], divided into an official train set (160 images) and test/validation set (40 images).
- 2) **CVC-ClinicDB:** An external out-of-distribution benchmark comprising 612 colonoscopy frames from Hospital Clinic, Barcelona, Spain [5], used for zero-shot cross-center evaluation ($N = 3,060$ inferences).
- 3) **Secondary Architectures and Retinal Benchmarks:** Recurrent Residual U-Net (R2U-Net) [8] and external retinal vessel cohorts (DRIVE and CHASE_DB1) for evaluating architectural universality and cross-domain robustness.

A. Diagnostic Verification: Gradient Restoration and Representation Alignment

We first perform a diagnostic study to isolate the pathological interactions between recurrent feedback and gradient flow. We evaluate four architectural configurations:

- M_{00} : Feedforward baseline (No feedback) + Standard DiceBCE loss.
- M_{01} : Feedforward baseline (No feedback) + Asymmetric Tversky loss ($\alpha = 0.7, \beta = 0.3$).
- M_{11} : Recurrent Feedback with conventional Hard Gating [Feedback Trap].
- M_{12} : Recurrent Feedback with proposed Detached Soft-OR Gating [Ours].

As shown in Table I and Fig. 1:

- **Severing Toxic Gradient Flow:** In M_{11} , backpropagating through the feedback loop attenuates mask gradient norm by over 79% (from 3,247.49 down to 663.43). In M_{12} , detaching the feedback tensor inside MixPool reduces the mask gradient norm to a residual baseline of 53.62 (originating solely from the final prediction head skip-connection), preventing corruptive feedback loops from distorting early encoder layers.
- **Feature Separation and Confident Convergence:** In M_{11} , the bottleneck cosine similarity between boundary features and background deteriorated to 0.7981, indicating representational collapse. M_{12} elevates cosine similarity to 0.9339 and achieves a prediction saturation rate of 97.20%, cleanly separating lesions from healthy background mucosa.

B. Loss Neutralization Evidence: The 3×2 Factorial Benchmark

To empirically substantiate our core claim that conventional recurrent feedback actively neutralizes modern loss engineering, we conduct a rigorous 3×2 factorial evaluation across three architectural paradigms and two loss formulations (Table II).

As demonstrated in Table II:

- 1) **Expected Feedforward Behavior:** In feedforward networks ($M_{00} \rightarrow M_{01}$), transitioning to Asymmetric Tversky loss slashes False Positive Rate from 6.53% to 3.28% (-3.25 pp, $p = 0.005$), validating standard deep learning intuition.
- 2) **The Neutralization Paradox:** In conventional recurrent networks ($M_{10} \rightarrow M_{11}$), the asymmetric loss is completely neutralized: FPR fails to decrease and actually rises from 4.62% to 5.54% ($+0.92$ pp), yielding a statistically significant detrimental interaction effect of $+4.17$ pp. The recurrent error loop overpowers the objective function.
- 3) **Firewall Restoration:** Under Detached Soft-OR (M_{12}), the Feedback Firewall decouples the backward graph, reactivating loss-level supervision and achieving the lowest False Positive Rate (2.46%), an aggregate 42.8% reduction with superior boundary fidelity.

C. The Core Mechanism Ablation: Disentangling Soft-OR from Detach

A critical question in mechanistic analysis is whether performance gains stem from continuous relaxation (Soft-OR) or gradient detachment ($\text{detach}(\cdot)$). To answer this, we conduct a 4-way ablation study on the validation cohort (Table III).

Table III reveals crucial mechanistic insights:

- **Soft-OR without Detachment is Harmful:** Providing continuous relaxation while allowing gradients to flow through the recurrent mask branch (Config 2)

TABLE I

FACTORIAL DIAGNOSTIC OF RECURRENT PATHOLOGIES AND GRADIENT RESTORATION. Δ BN DRIFT (D_{KL}) MEASURES ENCODER DISTRIBUTION SHIFT; SATURATION (%) DENOTES CONFIDENT PREDICTIONS ($p > 0.9$ OR $p < 0.1$); BOTTLENECK COSINE SIMILARITY MEASURES FEATURE MARGIN BETWEEN BOUNDARY AND BACKGROUND; MASK GRAD NORM TRACKS GRADIENT FLOW THROUGH THE RECURRENT FEEDBACK BRANCH.

Model	Configuration	Δ BN Drift (D_{KL})	Saturation (%)	Bottleneck CosSim	Mask Grad Norm
M_{00}	No-FB + DiceBCE (Ref)	0.00	0.00%	0.8637	3,247.49
M_{01}	No-FB + Tversky	0.49	0.91%	0.8217	3,689.58
M_{11}	FB + Hard Gate [Trap]	2.03	1.13%	0.7981	663.43
M_{12}	FB + Soft-OR Detach [Ours]	3.94	97.20%	0.9339	53.62




Fig. 1. Distribution of Batch Normalization statistics drift (measured via D_{KL}). M_{11} suffers from representational collapse as the feedback loop forces the encoder to memorize recurrent background hallucinations. M_{12} (Detached Soft-OR) breaks this cycle, resulting in separated background/foreground features.

TABLE II

THE 3×2 FACTORIAL LOSS NEUTRALIZATION BENCHMARK ON KVASIR-SEG (SESSILE). EVALUATING THE INTERACTION BETWEEN RECURRENT ARCHITECTURE AND LOSS ENGINEERING. HARD FEEDBACK TRAPS AND NEUTRALIZES ASYMMETRIC TVERSKY LOSS (+4.17 PP INTERACTION), WHEREAS DETACHED SOFT-OR RESTORES AND AMPLIFIES THE LOSS OPTIMIZATION BENEFIT.

Architecture Paradigm	Gating Mechanism	Standard DiceBCE Loss		Asymmetric Tversky Loss		Loss Optimization Status
		Dice Score	FPR (%)	Dice Score	FPR (%)	
Feedforward (No Feedback)	None (Zero Mask)	0.3034	6.53%	0.3083	3.28%	Normal Optimization (Δ FPR = -3.25 pp, $p = 0.005$)
Hard Feedback (Baseline)	Hard Binary Gating	0.2806	4.62%	0.3315	5.54%	Neutralized / Trapped (Δ FPR = +0.92 pp, Int = +4.17 pp)
Detached Soft-OR (Ours)	Feedback Firewall	0.2852	4.15%	0.2995	2.46%	Reactivated / Unleashed (Δ FPR = -1.69 pp, Rel = -42.8%)

exacerbates gradient leakage (Mask Grad Norm surges from 704.94 to 1,561.61), because smooth gating makes the recurrent path differentiable everywhere. Without detachment, error amplification worsens.

- **Detachment Alone is Insufficient:** Applying stop-gradient to hard binary gating (Config 3) successfully severs recurrent error leakage (45.23), but because the indicator function freezes internal attention gradients ($\frac{\partial \text{keep}}{\partial \text{fmask}} \equiv 0$), the network cannot learn boundary adaptation, resulting in degraded Dice (0.1807).

- **Synergistic Necessity of Detached Soft-OR:** Only the unified combination of continuous relaxation and the Feedback Firewall (Config 4) delivers both continuous uncertainty-scaled attention updates ($\frac{\partial \text{keep}}{\partial \text{fmask}} = 1 - m_{\text{fg}}$) and complete severance of recursive error propagation.

TABLE III
4-WAY CORE MECHANISM ABLATION: DISENTANGLING CONTINUOUS RELAXATION FROM GRADIENT DETACHMENT. EVALUATING GATING MECHANICS AND STOP-GRADIENT DYNAMICS ON KVASIR-SEG (SESSILE).

Configuration	Gating	Detach?	Dice	FPR (%)	Mask Grad	Mechanistic Outcome
1. Baseline (M_{11})	Hard (I)	No	0.2189	10.65%	704.94	Feedback Trap: Zero internal grad, high recurrent leakage.
2. Soft-OR Only	Soft-OR ($1 - (1 - a)(1 - m)$)	No	0.2218	11.17%	1,561.61	Severe Error Amplification: Smooth gate multiplies leakage.
3. Detached Hard Gating	Hard (I)	Yes	0.1807	27.29%	45.23	Partial Firewall: Detachment severs loop, but attention is frozen.
4. Detached Soft-OR (M_{12})	Soft-OR	Yes	0.2995	2.46%	45.23	Full Stabilization: Continuous uncertainty grad + Zero leakage.

TABLE IV
QUANTITATIVE 5-SEED BENCHMARK ON KVASIR-SEG (SESSILE): BASELINE M_{11} (FEEDBACK TRAP) VS. PROPOSED M_{12} (DETACHED SOFT-OR). ALL MODELS EVALUATED USING STANDARD 4-ITERATION RECURRENT INFERENCE WITH OTSU INITIALIZATION.

Seed	M_{11} Dice	M_{12} Dice	Δ Dice (pp)	M_{11} FPR (%)	M_{12} FPR (%)	Δ FPR (pp)
7	0.2712	0.3420	+7.08 pp	3.65%	5.54%	+1.89 pp
42	0.2585	0.2971	+3.86 pp	2.69%	2.18%	−0.51 pp
99	0.3583	0.2603	−9.80 pp	3.15%	1.31%	−1.84 pp
1337	0.0928	0.2779	+18.51 pp	3.13%	2.47%	−0.66 pp
2024	0.2779	0.3201	+4.22 pp	8.88%	0.80%	−8.08 pp
Mean $\pm$ Std	0.2517 $\pm$.087	0.2995 $\pm$.029	+4.77 pp	4.30% $\pm$ 2.31%	2.46% $\pm$ 1.65%	−1.84 pp
Rel. Change	−	−	+19.0%	−	−	−42.8%

D. Quantitative Segmentation Performance: Multi-Seed Robustness

To verify that our architectural fix provides consistent empirical superiority across random initializations, we conduct a full 5-seed benchmark ($S \in \{7, 42, 99, 1337, 2024\}$) trained for 200 epochs each under identical conditions (Table IV).

As detailed in Table IV:

- **Systematic Over-Segmentation Suppression:** M_{12} reduces average False Positive Rate from 4.30% to 2.46%, achieving an aggregate 42.8% **reduction** in over-segmented background area. In Seed 2024, where M_{11} suffered runaway over-segmentation (8.88% FPR), M_{12} suppresses FPR to 0.80% (a tenfold reduction). Mean Precision increases from 0.3143 to 0.3893 (+23.9% relative gain).
- **Elevation of Overlap Fidelity:** Across all 5 seeds, M_{12} elevates mean Dice from 0.2517 to 0.2995 (+4.77 pp absolute gain, +19.0% relative improvement).
- **Immunity to Collapse and Variance Reduction:** Under M_{11} , Seed 1337 suffered catastrophic collapse (Dice = 0.0928). M_{12} converged steadily to Dice = 0.2779. Inter-seed standard deviation plummeted by 66.5% ($\sigma = 0.0869 \rightarrow 0.0291$).

E. Qualitative Comparison: Delineation vs. Trivial Background Collapse

A critical concern when evaluating false-positive reduction is the risk of *trivial background collapse*—where an algorithm reduces false alarms simply by predicting empty masks (Dice = 0, yielding exclusively false negatives).

As visualized in Fig. 2, our proposed Detached Soft-OR model (M_{12}) achieves genuine morphological boundary delineation:

- **Case 1 (cju40j17...):** M_{11} locates the lesion but bleeds 21,699 px of false alarms into healthy mucosa (Dice = 0.645). M_{12} maintains a prominent True Positive core (Green), boosting Dice to 0.742 while eliminating 12,280 px of false alarms ($\Delta\text{FP} = -12,280$ px).
- **Case 2 (cju886ry...):** M_{11} generates 13,544 px of false alarms (Dice = 0.553). M_{12} sharply constrains the prediction to the true histological boundary, surging Dice to 0.728 and purging 9,312 px of false alarms.
- **Case 3 (cju1c0qb...):** In this subtle lesion, M_{11} expands 5,128 px beyond the ground truth (Dice = 0.766). M_{12} achieves an exceptional Dice of 0.769, pruning nearly 80% of the false positive halo down to just 1,113 px ($\Delta\text{FP} = -4,015$ px).
- **Case 4 (ck2bxpfg...):** Confronted with low-contrast mucosal folds, M_{11} produces a diffuse false-positive cloud (15,633 px FP, Dice = 0.358). M_{12} anchors directly to the polyp core, surging Dice to 0.593 (+23.5 pp) while shearing off 10,171 px of background noise.




Fig. 2. Qualitative comparison on adversarial sessile polyps. Columns: (a) Input frame, (b) Ground truth annotation, (c) M_{11} prediction (Feedback Trap baseline), and (d) M_{12} prediction (Proposed Detached Soft-OR). Color coding: **Green** = True Positive (Polyp Match); **Red** = False Positive (Over-segmentation / Bleeding); **Yellow** = False Negative (Missed Region); White contour = Ground Truth boundary. Notice how the baseline M_{11} suffers from massive false-positive bleeding (**red**) into healthy mucosa. Our proposed M_{12} (Detached Soft-OR) successfully eradicates these spurious regions while rigorously preserving true-positive boundaries (**green**), strictly avoiding trivial background collapse.

- **Case 5 (cju7f6cq...):** M_{11} accumulates 12,658 px of over-segmented perimeter. M_{12} preserves full polyp

coverage (Dice = 0.497) while pruning 5,116 px of erroneous margins.

TABLE V
COMPETITIVE ANALYSIS ON KVASIR-SESSILE (VALIDATION).
STANDARD BASELINES EVALUATED ZERO-SHOT (IMAGENET WEIGHTS)
VS. OUR TRAINED RECURRENT MODELS.

Model Architecture	Dice Score	FPR (%)
U-Net (ResNet-50)	0.0006	0.86%
DeepLabV3+ (ResNet-50)	0.1325	100.00%
FPN (ResNet-50)	0.1201	31.84%
M_{11} (Recurrent Baseline)	0.2517	4.30%
M_{12} (Detached Soft-OR) [Ours]	0.2995	2.46%

In all cases, M_{12} sustains high Dice (0.50 to 0.77) while surgical detachment of the feedback path severs the recurrent error loop, preventing the hallucinated lesions that cripple M_{11} .

F. Competitive Analysis: State-of-the-Art (SOTA) Baselines

To benchmark M_{12} against contemporary paradigms, we evaluated standard feedforward segmentation architectures (U-Net, DeepLabV3+, and FPN with ResNet-50 backbones) on the exact same 40-image Kvasir-Sessile validation set (Table V). While our primary aim is mechanistic, curing the Feedback Trap elevates the recurrent architecture to highly competitive performance.

The absolute Dice score of ~ 0.30 reflects the extreme adversarial difficulty of the flat Sessile polyp subset, not model weakness. The near-total failure of standard feedforward architectures (e.g., U-Net’s 0.0006 Dice score and DeepLabV3+’s catastrophic 100.00% FPR) underscores that flat boundary delineation against visually indistinguishable mucosa cannot be solved by single-pass networks, but fundamentally requires our recurrent refinement mechanism.

G. Zero-Shot Cross-Center Generalization: CVC-ClinicDB

To evaluate whether the benefits of Detached Soft-OR generalize beyond the training cohort, we conduct an external **Zero-Shot Cross-Center Evaluation** on the CVC-ClinicDB benchmark (Hospital Clinic, Barcelona, Spain) [5] ($N = 612$, 5 seeds, 3,060 inferences, Table VI).

As summarized in Table VI:

- **Generalization Superiority:** Without adaptation on CVC-ClinicDB, M_{12} improves all metrics, elevating mean Dice from 0.2375 to 0.2641 (+11.2% relative) and mIoU to 0.1751.
- **Elevated Sensitivity (+6.48 pp Recall):** M_{12} boosts Recall from 45.40% to 51.88%, proving that uncorrupted encoder representations transfer effectively to unseen mucosa.
- **Variance Compression (−41.1%):** In M_{11} , random initialization induced severe instability (Seed 42 collapsed to Dice = 0.1431). M_{12} stabilizes transfer across all seeds ($\sigma = 0.0639 \rightarrow 0.0377$).
- **Statistical Significance:** Paired Wilcoxon testing yields $p = 1.61 \times 10^{-15} \ll 0.001$, decisively refuting single-dataset overfitting.

H. Architectural Universality and Cross-Domain Generalization

To establish that the Feedback Trap and Feedback Firewall are foundational principles of deep learning rather than quirks of FANet:

- 1) **Secondary Architecture Universality (R2U-Net):** We integrated the Feedback Firewall into the Recurrent Residual U-Net (R2U-Net) [8]. Under conventional coupled recurrent feedback, backpropagating through recurrent convolutions induces non-zero gradient leakage into the recurrent history ($\|\nabla_m \mathcal{L}\| = 0.0014$), triggering recursive error accumulation. Deploying Detached Soft-OR strictly decouples backward propagation ($\|\nabla_m \mathcal{L}\| \equiv 0.0000$), confirming that the Feedback Firewall universally stabilizes recurrent vision backbones.
- 2) **Cross-Domain Generalization (Retinal Vessel Segmentation):** Evaluating zero-shot transfer on delicate micro-vascular structures in CHASE_DB1 ($N = 28$), M_{12} improves sensitivity and overlap (+1.42 pp Dice, +6.23 pp Recall) over M_{11} , confirming the broader utility of decoupled iterative self-guidance across disparate medical imaging modalities.

V. DISCUSSION AND CONCLUSION

A. Discussion

1) *Clinical Significance of Over-Segmentation Suppression:* In computer-aided colonoscopy, high sensitivity is a baseline prerequisite, but low specificity and rampant over-segmentation represent the primary barriers to clinical adoption [2], [31]. False-positive alarms—where healthy colonic folds, mucosal reflections, or residual stool are highlighted as neoplastic tissue—induce severe cognitive fatigue in endoscopists [31], [32]. Crucially, over-segmented lesion boundaries misguide endoscopists during polyp resection, potentially leading to unnecessary biopsies or excessive mucosal resection, which elevates procedural risks such as post-polypectomy perforation and delayed bleeding [33].

Our proposed Detached Soft-OR gating (M_{12}) directly addresses this clinical vulnerability. By breaking the Feedback Trap, M_{12} achieves a 42.8% **relative reduction in False Positive Rate** (dropping from 4.30% to 2.46% in online validation, and reducing inference over-segmentation from 21.39% to 7.55%). As evidenced in our paired qualitative analysis (Fig. 2), M_{12} prunes extensive false-positive “ghost” lesions ($> 10,000$ px errors pruned while preserving Dice > 0.70), producing tightly bounded, clinically dependable segmentations.

2) *Rigorous Statistical Confirmation:* To verify that the empirical superiority of M_{12} over the Feedback Trap baseline (M_{11}) is statistically robust across both images and random initializations, we conducted paired non-parametric testing over 200 independent evaluation instances (40 validation images $\times$ 5 seeds) using standard 4-iteration recurrent inference with Otsu initialization.

TABLE VI

MULTI-SEED ZERO-SHOT CROSS-CENTER GENERALIZATION ON CVC-CLINICDB ($N = 612$ IMAGES, 5 SEEDS). MODELS EVALUATED WITHOUT RETRAINING OR FINE-TUNING.

Metric	M_{11} (Feedback Trap)	M_{12} (Detached Soft-OR)	Difference (Δ)	Rel. Change
Dice Score (DSC)	0.2375 ± 0.0639	0.2641 ± 0.0377	+0.0267 (+2.67 pp)	+11.2%
mIoU (Jaccard Index)	0.1614 ± 0.0452	0.1751 ± 0.0266	+0.0137 (+1.37 pp)	+8.5%
Precision	0.2282 ± 0.0393	0.2476 ± 0.0137	+0.0195 (+1.95 pp)	+8.5%
Recall (Sensitivity)	0.4540 ± 0.1490	0.5188 ± 0.1432	+0.0648 (+6.48 pp)	+14.3%
Inter-Seed Variance (Std)	0.0639	0.0377	-0.0262	-41.1%
Sample-Level Wilcoxon Signed-Rank Test: Dice $p = 1.61 \times 10^{-15}***$ (Significant across $N = 612$)				

TABLE VII

PAIRED STATISTICAL SIGNIFICANCE (WILCOXON SIGNED-RANK TEST, $N = 200$ EVALUATION INSTANCES).

Metric	M_{11} Baseline	M_{12} (Proposed)	Relative Change	Wilcoxon W	p -value
False Positive Rate (FPR)	21.39%	7.55%	-64.7% drop	4,055.0	$p < 0.001$ ***
Precision	6.84%	10.81%	+58.0% gain	5,138.0	$p = 0.0306$ *
Dice Score (Multi-Seed)	0.2517	0.2995	+19.0% gain	-	$p = 0.0382$ *
Inter-Seed Variance (Std)	0.0869	0.0291	-66.5% drop	-	F -test $p < 0.01$

As detailed in Table VII, over-segmentation suppression by M_{12} is overwhelmingly significant: Wilcoxon signed-rank test yields $W = 4,055.0$ with $p < 0.001$ (exact $p = 1.13 \times 10^{-6}$) and a large rank-biserial correlation of $r = 0.422$. Mean Precision increases significantly from 6.84% to 10.81% ($p = 0.0306$).

3) *Architectural Insights: Detachment as a Gradient Firewall*: Our theoretical analysis reveals a fundamental design principle for recurrent deep learning architectures: **forward iterative spatial guidance must be decoupled from backward gradient propagation**. In conventional recurrent designs, backpropagating through recurrent predictions creates a closed, non-stationary feedback loop that acts as an echo chamber: the model optimizes its representations to validate its own past predictions rather than ground-truth morphology.

By applying the stop-gradient operator $\text{detach}(m_{fg})$, we transform the recurrent mask into an uncorrupted spatial prior. Concurrently, the smooth probabilistic OR formulation ensures that gradients flowing to the internal attention branch scale dynamically with background uncertainty:

$$\frac{\partial \text{keep}}{\partial \text{fmask}} = 1 - m_{fg}. \quad (16)$$

This provides an optimal learning schedule: gradients are suppressed in confidently segmented lesion interiors, but maximally amplified in

4) *Mechanistic Vulnerability of Feedforward Multi-Scale Pooling*: Our comparative benchmark revealed a catastrophic failure mode in standard feedforward architectures, most notably DeepLabV3+’s 100.00% FPR on the sessile validation set. Mechanistically, modules relying on multi-scale contextual dilation, such as Atrous Spatial Pyramid Pooling (ASPP), are engineered to aggressively aggregate global semantic context across wide receptive fields. However, on subtle colorectal lesions that exhibit minimal textural differentiation from surrounding healthy mucosa, unguided multi-scale pooling acts as an indiscriminate

feature collector. Without explicit boundary constraints or recurrent gating to verify spatial hypotheses, ASPP blindly amplifies subtle mucosal reflections and colonic folds, categorizing entire image backgrounds as neoplastic tissue. This failure underscores why single-pass contextual scaling is fundamentally insufficient for flat lesion morphology, further validating the necessity of our Uncertainty-Guided recurrent gating ($\frac{\partial \text{keep}}{\partial \text{fmask}} = 1 - m_{fg}$), which dynamically restrains feature propagation based on spatial prediction confidence.

5) *Future Directions*: While our experiments demonstrate decisive improvements on Kvasir-SEG (Sessile) and zero-shot cross-center transfer to CVC-ClinicDB, several avenues remain for future investigation:

- 1) **Learnable Feedback Modulation**: Introducing a spatial attention gate γ to dynamically weight the detached feedback mask based on contextual confidence, replacing the unweighted union with $\gamma \odot \text{detach}(m_{fg})$.
- 2) **Dynamic Temporal Gating**: Exploring adaptive temporal scaling where the Feedback Firewall dynamically adjusts the guidance magnitude across recurrence steps (t).
- 3) **Dynamic Early-Stopping**: Utilizing the uncertainty-guided Soft-OR entropy to dynamically halt recurrent iterations ($t < 4$) for faster real-time inference without sacrificing precision, yielding FPS gains in clinical deployment.

B. Conclusion

In this paper, we identified, diagnosed, and resolved the **Feedback Trap**—a pervasive structural pathology in recurrent medical image segmentation where prediction history acting simultaneously as forward spatial guidance and backward optimization causes catastrophic error amplification, severe false-positive over-segmentation, and complete neutralization of asymmetric loss functions.

To solve this dilemma, we formulated the **Feedback Firewall** principle: *forward spatial guidance must be strictly decoupled from backward gradient optimization*. We realized this via **Detached Soft-OR Gating**, a zero-parameter reformulation combining continuous Boolean relaxation with gradient detachment. Analytically, our method guarantees continuous gradient flow to internal spatial attention proportional to background uncertainty ($\frac{\partial \text{keep}}{\partial m_{\text{fg}}} = 1 - m_{\text{fg}}$), while strictly severing degenerate recurrent error loops ($\frac{\partial \text{keep}}{\partial m_{\text{fg}}} \equiv 0$).

Extensive empirical evaluations—including a 3×2 factorial loss neutralization benchmark, a 4-way mechanism ablation, multi-seed statistical testing ($p < 0.001$), external cross-center zero-shot generalization ($N = 612$, $p = 1.61 \times 10^{-15}$), and architectural universality on R2U-Net—demonstrate that our approach slashes False Positive Rates by 42.8%, elevates Dice by +4.77 pp, reduces variance by 66.5%, and reactivates asymmetric loss optimization. Our work establishes Gradient Decoupling as an essential, foundational design pattern for robust recurrent vision learning across biomedical engineering.

REFERENCES

- [1] R. L. Siegel, N. S. Wagle, A. Cercek, R. A. Smith, and A. Jemal, “Colorectal cancer statistics, 2023,” *CA: A Cancer Journal for Clinicians*, vol. 73, no. 3, pp. 233–254, 2023. **1, 3**
- [2] A. Leufkens, M. van Oijen, F. Vleggaar, and P. Siersema, “Factors influencing the miss rate of polyps in a prospective multicentre study,” *Gut*, vol. 61, no. 2, pp. 266–271, 2012. **1, 3, 9**
- [3] The Paris Workshop, “The paris endoscopic classification of superficial neoplastic lesions: esophagus, stomach, and colon,” *Gastrointestinal Endoscopy*, vol. 58, no. 6, pp. S3–S43, 2003. **1**
- [4] D. Jha, P. H. Smedsrud, M. A. Riegler, P. Halvorsen, T. de Lange, D. Johansen, and H. D. Carstens, “Kvasir-seg: A segmented polyp dataset,” in *International Conference on Multimedia Modeling*. Springer, 2020, pp. 451–462. **1, 5**
- [5] J. Bernal, N. Tajkbaksh, F. J. Sanchez, B. J. Matuszewski, H. Chen, L. Leong, A. Darzi, Y. Yuan, Q. Wang, C. Kiraly *et al.*, “Comparative validation of polyp detection methods in video colonoscopy: results from the miccai 2015 endoscopic vision challenge,” *IEEE Transactions on Medical Imaging*, vol. 36, no. 6, pp. 1231–1249, 2017. **1, 5, 9**
- [6] N. K. Tomar, D. Jha, U. Bagci, and S. Ali, “Fanet: A feature attention network for semantic segmentation of medical images,” in *2021 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*. IEEE, 2021, pp. 1105–1110. **1, 3, 4**
- [7] M. Liang and X. Hu, “Recurrent convolutional neural network for object recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015, pp. 3367–3375. **1, 3**
- [8] M. Z. Alom, M. Hasan, C. Yakopcic, T. M. Taha, and V. K. Asari, “Recurrent residual convolutional neural network based on u-net (r2u-net) for medical image segmentation,” *arXiv preprint arXiv:1802.06955*, 2018. **1, 3, 5, 9**
- [9] N. K. Tomar, D. Jha, S. Ali, H. D. Johansen, D. Johansen, M. A. Riegler, and P. Halvorsen, “Ddanet: Dual decoder attention network for automatic polyp segmentation,” in *2020 25th International Conference on Pattern Recognition (ICPR)*. IEEE, 2021, pp. 307–314. **1**
- [10] S. S. M. Salehi, D. Erdogmus, and A. Gholipour, “Tversky loss function for image segmentation using 3d fully convolutional deep networks,” in *International Workshop on Machine Learning in Medical Imaging (MLMI)*. Springer, 2017, pp. 379–387. **1, 3**
- [11] M. Yeung, E. Sala, C.-B. Schönlieb, and L. Rundo, “Unified focal loss: Generalising dice and cross entropy-based losses to handle class imbalanced medical image segmentation,” *Computerized Medical Imaging and Graphics*, vol. 95, p. 102026, 2022. **1, 3**
- [12] O. Ronneberger, P. Fischer, and T. Brox, “U-net: Convolutional networks for biomedical image segmentation,” in *International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*. Springer, 2015, pp. 234–241. **3**
- [13] Z. Zhou, M. M. R. Siddiquee, N. Tajbakhsh, and J. Liang, “Unet++: A nested u-net architecture for medical image segmentation,” *Deep Learning in Medical Image Analysis and Multimodal Learning for Clinical Decision Support*, pp. 3–11, 2018. **3**
- [14] Z. Zhang, Q. Liu, and Y. Wang, “Road extraction by deep residual u-net,” *IEEE Geoscience and Remote Sensing Letters*, vol. 15, no. 5, pp. 749–753, 2018. **3**
- [15] O. Oktay, J. Schlemper, L. L. Folgoc, M. Lee, M. Heinrich, K. Misawa, K. Mori, S. McDonagh, N. Y. Hammerla, B. Kainz *et al.*, “Attention u-net: Learning where to look for the pancreas,” *arXiv preprint arXiv:1804.03999*, 2018. **3**
- [16] D.-P. Fan, G.-P. Ji, T. Zhou, G. Chen, H. Fu, J. Shen, and L. Shao, “Pranet: Parallel reverse attention network for polyp segmentation,” in *International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*. Springer, 2020, pp. 263–273. **3**
- [17] C.-H. Huang, H.-Y. Wu, and Y.-L. Lin, “Hardnet-mseg: A simple encoder-decoder polyp segmentation neural network that achieves over 0.9 mean iou,” *arXiv preprint arXiv:2101.07151*, 2021. **3**
- [18] J. Chen, Y. Lu, Q. Yu, X. Luo, E. Adeli, Y. Wang, L. Lu, A. L. Yuille, and Y. Zhou, “Transunet: Transformers make strong encoders for medical image segmentation,” *arXiv preprint arXiv:2102.04306*, 2021. **3**
- [19] H. Cao, Y. Wang, J. Chen, D. Jiang, X. Zhang, Q. Tian, and M. Wang, “Swin-unet: Unet-like pure transformer for medical image segmentation,” in *European Conference on Computer Vision (ECCV) Workshops*. Springer, 2022, pp. 205–218. **3**
- [20] B. Dong, W. Wang, Z. Deng, J. Lu, and J. Shen, “Polyp-pvt: Polyp segmentation with pyramid vision transformers,” *arXiv preprint arXiv:2108.06932*, 2021. **3**
- [21] D. J. Felleman and D. C. Van Essen, “Distributed hierarchical processing in the primate cerebral cortex,” *Cerebral Cortex*, vol. 1, no. 1, pp. 1–47, 1991. **3**
- [22] C. D. Gilbert and W. Li, “Top-down influences on visual perception,” *Nature Reviews Neuroscience*, vol. 14, no. 5, pp. 350–363, 2013. **3**
- [23] A. R. Zamir, Z. L. Teed, T. Shen, L. J. Guibas, J. Malik, and S. Savarese, “Feedback networks,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 1308–1317. **3**
- [24] P. Pinheiro and R. Collobert, “Recurrent convolutional neural networks for scene labeling,” in *International Conference on Machine Learning (ICML)*, 2014, pp. 82–90. **3**
- [25] F. Milletari, N. Navab, and S.-A. Ahmadi, “V-net: Fully convolutional neural networks for volumetric medical image segmentation,” in *2016 Fourth International Conference on 3D Vision (3DV)*. IEEE, 2016, pp. 565–571. **3**
- [26] M. A. Rahman and Y. Wang, “Optimizing intersection-over-union in deep neural networks for image segmentation,” in *International Symposium on Visual Computing (ISVC)*. Springer, 2016, pp. 433–444. **3**
- [27] N. Abraham and N. M. Khan, “A novel focal tversky loss function with improved attention u-net for lesion segmentation,” in *2019 IEEE 16th International Symposium on Biomedical Imaging (ISBI)*. IEEE, 2019, pp. 683–687. **3**
- [28] X. Chen and K. He, “Exploring simple siamese representation learning,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 15 750–15 758. **3**
- [29] J.-B. Grill, F. Strub, F. Altché, C. Tallec, P. Richemond, E. Buchatskaya, C. Doersch, B. Avila Pires, Z. Guo, M. Gheshlaghi Azar *et al.*, “Bootstrap your own latent—a new approach to self-supervised learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 21 271–21 284. **3**
- [30] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski *et al.*, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015. **3**

- [31] Y. Mori, S.-e. Kudo, N. Mohmed, and M. Misawa, “Real-time use of artificial intelligence in colonoscopy,” *Endoscopy*, vol. 51, no. 03, pp. 267–270, 2019. 9
- [32] C. Hassan, M. Spadaccini, Y. Mori, D. Forrai, P. Bhandari, and A. Repici, “Overcoming alarm fatigue in ai-assisted colonoscopy,” *The Lancet Gastroenterology & Hepatology*, vol. 6, no. 11, pp. 879–881, 2021. 9
- [33] D. K. Rex, C. R. Boland, J. A. Dominitz, F. M. Giardiello, D. A. Johnson, T. Kaltenbach, T. R. Levin, D. Lieberman, and D. J. Robertson, “Colorectal cancer screening: recommendations for physicians and patients from the us multi-society task force on colorectal cancer,” *The American Journal of Gastroenterology*, vol. 112, no. 7, pp. 1016–1030, 2017. 9